

Fast Layout-Oblivious Tensor-Matrix Multiplication with BLAS

Cem Savaş Başsoy

Hamburg University of Technology, Schwarzenbergstrasse 95, Germany,
cem.bassoy@gmail.com

Abstract. The tensor-matrix product is a basic tensor operation that is required by various tensor methods such as the ALS or the HOSVD. This paper presents flexible high-performance algorithms for the mode- q tensor-matrix multiplication that computes the product according to the the Loops-over-gemms (LOG) approach with `gemm` being the general matrix multiply. Our algorithms can process dense tensors with any linear tensor layout, arbitrary tensor order and dimensions all of which can be runtime variable. We discuss different tensor slicing methods with parallelization strategies and propose six variations of a base algorithm which calls a `gemv`, `gemm` and/or a `gemm_batch` with subtensors or tensor slices. Their performance is quantified for a set of tensors with various shapes and tensor order. The best-performing version attains a median performance of 1.37 double precision Tflops/s on an Intel Xeon Gold 6248R processor using Intel’s MKL. We show that the performance is only slightly affected by the tensor layout and the median performance is between [?] and [?] Tflops/s for a range of linear tensor formats. Our fastest version of the tensor-matrix multiplication is on average at least 14.05% and up to 3.79 x faster than other state-of-the-art implementations, including Libtorch and Eigen.

1 Introduction

Tensor computations are found in many scientific fields such as computational neuroscience, pattern recognition, signal processing and data mining [7, 14]. Tensors representing large amount of multidimensional data are decomposed and analyzed with the help of basic tensor operations [8, 9]. The decomposition and analysis led to the development and analysis of high-performance kernels for tensor contractions. In this work, we present and analyze a high-performance algorithm for the tensor-matrix multiplication that is used in many numerical algorithms such as the alternating least squares method [8, 9]. It is a compute-bound tensor operation and has the same arithmetic intensity as a matrix-matrix multiplication which can almost reach the practical peak performance of a computing machine.

There has been three main approaches for implementing tensor contractions. The Transpose-Transpose-gemm-Transpose (TGGT) approach reorganizes (flattens) tensors in order to perform a tensor contraction with an optimized

matrix-matrix multiplication (`gemm`) implementation [2, 16]. Implementations of a more recent method (GETT) are based on high-performance `gemm`-like algorithms [1, 12, 17]. A different method is the LOG approach in which BLAS are utilized with multiple tensor slices or subtensors if possible [10, 13, 15]. Implementations of the LOG and TTGT approaches are in general easier to maintain and faster to port than GETT implementations which might need to adapt vector instructions or blocking parameters according to a processor’s micro-architecture.

Our work is motivated by the fact that LOG-based implementations of the tensor-matrix and tensor-vector multiplication are similar. To our best knowledge, we are the first to combine the approach in [3] with the findings in [10] and to propose fast in-place tensor-matrix multiplication algorithms that are layout-oblivious. Our algorithms compute the tensor-matrix product in parallel using OpenMP together with highly efficient `gemm` or `gemm_batch` implementations. They support dense tensors with any order, dimensions and any linear tensor layout including the first- and the last-order storage formats for any contraction mode all of which can be runtime variable. Input and output tensors do not need to be transposed or flattened into two-dimensional matrices. The parallel versions of the recursive base algorithm execute fused loops in parallel and are able to fully utilize a processors compute units. Despite, the generality of our approach, every proposed algorithm can be implemented with less than 100 lines of C++ code where the complexity is hidden by the BLAS implementation and the corresponding selection of subtensors or tensor slices. We have provided an open and free reference C++ implementation of all algorithms and a python interface for convenience. While we have used Intel’s MKL for our benchmarks, the user is free to choose any other library that provides the BLAS interface.

The following analysis quantifies the impact of the tensor layout, the tensor slicing method and parallel execution of slice-matrix multiplications with varying contraction modes. The runtime measurements of our implementations are compared with those presented in [12, 17] including Libtorch and Eigen. In summary, the main findings of our work are:

- A tensor-matrix multiplication can be implemented by an in-place algorithm with 1 `gemv` and 7 `gemm` calls supporting all combinations of contraction mode, tensor order and dimensions for any linear tensor layout.
- Our algorithm with variable loop fusion and parallel slice-matrix multiplications is on average 17% faster than a single batched `gemm` call when the contraction and leading dimensions are greater than 256.
- All our proposed algorithms are layout oblivious and achieve at least a median throughput of [?] for any linear tensor layout.
- Our LOG-based tensor-times-matrix implementation are in general faster than TTGT- and GETT-based implementations that have been described in [12, 17] including actively developed libraries such as Libtorch and Eigen. Using symmetrically shaped tensors, an average speedup of [?] x to [?] x for single and double precision floating point computations can be achieved.

The remainder of the paper is organized as follows. Section 2 presents related work. Section 3 introduces the terminology used in this paper and defines the

85 tensor-vector multiplication. Algorithm design and methods for parallel execu-
 86 tion is discussed in Section 4. Section 5 describes the test setup and discusses
 87 the benchmark results in Section 6. Conclusions are drawn in Section 7.

88 2 Related Work

89 The authors in [13] discuss the efficient tensor contractions with highly optimized
 90 BLAS. Based on the LOG approach, they define requirements for the use of GEMM for
 91 class 3 tensor contractions and provide slicing techniques for tensors. The slicing
 92 recipe for the class 2 categorized tensor contractions contains a short description
 93 with a rule of thumb for maximizing performance. Runtime measurements cover
 94 class 3 tensor contractions.

95 The work in [10] presents a framework that generates in-place tensor-matrix
 96 multiplication according to the LOG approach. The authors present two strategies
 97 for efficiently computing the tensor contraction applying GEMMs with tensors.
 98 They report a speedup of up to 4x over the TTGT-based MATLAB tensor toolbox
 99 library discussed in [2]. Although many aspects are similar to our work, the
 100 authors emphasize the code generation of tensor-matrix multiplications using
 101 high-performance GEMM's.

102 The authors of [17] present a tensor-contraction generator TCCG and the GETT
 103 approach for dense tensor contractions that is inspired from the design of a high-
 104 performance GEMM. Their unified code generator selects implementations from
 105 generated GETT, LoG and TTGT candidates. Their findings show that among 48 dif-
 106 ferent contractions 15% of LoG based implementations are the fastest. However,
 107 their tests do not include the tensor-vector multiplication where the contraction
 108 exhibits at least one free tensor index.

109 Using also the GETT approach, the author presents in [12] a runtime flexible
 110 tensor contraction library. He describes block-scatter-matrix algorithm which
 111 uses a special layout for the tensor contraction. The proposed algorithm yields
 112 results that feature a similar runtime behavior to those presented in [17].

113 3 Background

114 **Notation** An order- p tensor is a p -dimensional array [11] where tensor ele-
 115 ments are contiguously stored in memory. We write a , \mathbf{a} , \mathbf{A} and $\underline{\mathbf{A}}$ in order to
 116 denote scalars, vectors, matrices and tensors. If not otherwise mentioned, we
 117 assume $\underline{\mathbf{A}}$ to have a tensor order that is greater than 2. The p -tuple \mathbf{n} with
 118 $\mathbf{n} = (n_1, n_2, \dots, n_p)$ will be referred to as a dimension tuple with $n_r > 1$. We
 119 will use round brackets $\underline{\mathbf{A}}(i_1, i_2, \dots, i_p)$ or $\underline{\mathbf{A}}(\mathbf{i})$ to denote a tensor element where
 120 $\mathbf{i} = (i_1, i_2, \dots, i_p)$ is a multi-index. A subtensor is denoted by $\underline{\mathbf{A}}'$ and references
 121 elements of a tensor $\underline{\mathbf{A}}$. They are specified with p index ranges and form a selec-
 122 tion grid. In this work, the index range shall either address all indices of a given
 123 mode or a single element that are given by single indices i_r with $1 \leq r \leq p$.
 124 Elements n'_r of a subtensor's dimension tuple \mathbf{n}' are therefore n_r if all indices
 125 of mode r are selected and 1 otherwise. We will annotate subtensors using only

their non-unit modes such as $\underline{\mathbf{A}}'_{u,v,w}$ where $n_u > 1, n_v > 1$ and $n_w > 1$ and $1 \leq u \neq v \neq w \leq p$. It is sufficient to only provide non-unit modes as the remaining single indices correspond to the loop induction variables of the following algorithms. A subtensor is called a slice $\underline{\mathbf{A}}'_{u,v}$ if the full range selection of $\underline{\mathbf{A}}$ occurs with only two modes. A fiber $\underline{\mathbf{A}}'_u$ is a tensor slice with only one dimension greater than 1.

Linear Tensor Layouts We use a layout tuple $\boldsymbol{\pi} \in \mathbb{N}^p$ to encode all linear tensor layouts including the first-order or last-order layout. They contain permuted tensor modes whose priority is given by their index. For instance, the first- and last-order storage formats are given by $\boldsymbol{\pi}_F = (1, 2, \dots, p)$ and $\boldsymbol{\pi}_L = (p, p-1, \dots, 1)$. An inverse layout tuple $\boldsymbol{\pi}^{-1}$ is defined by $\boldsymbol{\pi}^{-1}(\boldsymbol{\pi}(k)) = k$. Given a layout tuple $\boldsymbol{\pi}$ with p modes, the π_r -th element of a stride tuple is given by $w_{\pi_r} = \prod_{k=1}^{r-1} n_{\pi_k}$ for $1 < r \leq p$ and $w_{\pi_1} = 1$. Tensor elements of the π_1 -th mode are contiguously stored in memory. The location of tensor elements is determined by the tensor layout and the layout function. For a given tensor layout and stride tuple, a layout function $\lambda_{\mathbf{w}}$ maps a multi-index to a scalar index with $\lambda_{\mathbf{w}}(\mathbf{i}) = \sum_{r=1}^p w_r(i_r - 1)$. With $j = \lambda_{\mathbf{w}}(\mathbf{i})$ being the relative memory position of an element with a multi-index \mathbf{i} , reading from and writing to memory is accomplished with j and the first element's address of $\underline{\mathbf{A}}$.

Non-Modifying Flattening and Reshaping The flattening operation $\varphi_{r,q}$ transforms an order- p tensor $\underline{\mathbf{A}}$ to another order- p' view $\underline{\mathbf{B}}$ that has different a shape \mathbf{m} and layout $\boldsymbol{\tau}$ tuple of length p' with $p' = p - q + r$ and $1 \leq r < q \leq p$. It is related to the tensor unfolding operation as defined in [8, p.459] but neither changes the element ordering nor copies tensor elements. Given a layout tuple $\boldsymbol{\pi}$ of $\underline{\mathbf{A}}$, the flattening operation $\varphi_{r,q}$ is defined for contiguous modes $\hat{\boldsymbol{\pi}} = (\pi_r, \pi_{r+1}, \dots, \pi_q)$ of $\boldsymbol{\pi}$. Let $j = 0$ if $k \leq r$ and $j = q - r$ otherwise for $1 \leq k \leq p'$. Then the resulting layout tuple $\boldsymbol{\tau} = (\tau_1, \dots, \tau_{p'})$ of $\underline{\mathbf{B}}$ is given by $\tau_r = \min(\boldsymbol{\pi}_{r,q})$ and $\tau_k = \pi_{k+j} + s_k$ if $k \neq r$ where $s_k = |\{\pi_i \mid \pi_{k+j} > \pi_i \wedge \pi_i \neq \min(\hat{\boldsymbol{\pi}}) \wedge r \leq i \leq p\}|$. Elements of the corresponding shape tuple \mathbf{m} are given by $m_{\tau_r} = \prod_{k=r}^q n_{\pi_k}$ and $m_{\tau_k} = n_{\pi_{k+j}}$ if $k \neq r$.

The reshaping operation ρ transforms an order- p tensor $\underline{\mathbf{A}}$ to another order- p tensor $\underline{\mathbf{B}}$ with different shape \mathbf{m} and layout $\boldsymbol{\tau}$ tuples of length p . In this work, it permutes the shape and layout tuple simultaneously without changing the element ordering and without copying tensor elements. The operation ρ uses a permutation tuple $\boldsymbol{\rho} = (\rho_1, \dots, \rho_p)$ to only modify shape and layout tuples. Elements of the resulting shape tuple \mathbf{m} and the layout tuple $\boldsymbol{\tau}$ are given by $m_r = n_{\rho_r}$ and $\tau_r = \pi_{\rho_r}$, respectively.

Tensor-Matrix Multiplication (TTM) Let $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ be order- p tensors with shapes $\mathbf{n}_a = (n_1, \dots, n_q, \dots, n_p)$ and $\mathbf{n}_c = (n_1, \dots, n_{q-1}, m, n_{q+1}, \dots, n_p)$. Let \mathbf{B} be a matrix of shape $\mathbf{n}_b = (m, n_q)$. A mode- q TTM is denoted by $\underline{\mathbf{C}} = \underline{\mathbf{A}} \times_q \mathbf{B}$ where an element of $\underline{\mathbf{C}}$ is given by

$$\underline{\mathbf{C}}(i_1, \dots, i_{q-1}, j, i_{q+1}, \dots, i_p) = \sum_{i_q=1}^{n_q} \underline{\mathbf{A}}(i_1, \dots, i_q, \dots, i_p) \cdot \mathbf{B}(j, i_q) \quad (1)$$

with $1 \leq i_r \leq n_r$ and $1 \leq j \leq m$. The mode q is the *contraction mode* of the TTM with $1 \leq q \leq p$. The tensor-matrix multiplication generalizes the computational aspect of the two-dimensional case $\mathbf{C} = \mathbf{B} \cdot \mathbf{A}$ if $p = 2$ and $q = 1$. Its arithmetic intensity is equal to that of a matrix-matrix multiplication and is not memory-bound. In the following, we assume that the tensors $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ have the same tensor layout $\boldsymbol{\pi}$. Elements of matrix $\underline{\mathbf{B}}$ can be stored in either the column-major or row-major format.

4 Algorithm Design

4.1 Sequential Algorithm

The sequential baseline algorithm for Eq. 1 can be implemented with a single C++ function that supports tensors with arbitrary order, dimensions and any linear tensor layout. It consists of nested recursion with a control flow that is akin to algorithm 1 in [4] consisting of two **if** statements with an **else** branch. The body of the first **if** statement contains a recursive call that skips the iteration over the dimension n_q when $r = \hat{q}$ with $\pi_r = q$ and $\hat{q} = \pi_q^{-1}$ where π^{-1} is the inverse layout tuple. The second **if** statement contains multiple recursive calls for the modes $1 \leq r \neq \hat{q} \leq p$ with different multi-indices. The **else** branch is the base case and consists of two loops that compute a fiber-matrix product. The outer loop iterates with j over the dimension m of $\underline{\mathbf{C}}$ and $\underline{\mathbf{B}}$. The inner loop iterates with i_q over the dimension n_q of $\underline{\mathbf{A}}$ and $\underline{\mathbf{B}}$ computing an inner product.

4.2 Baseline Algorithm with Contiguous Memory Access

The baseline algorithm accesses elements of $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ non-contiguously whenever $\pi_1 \neq q$. Matrix $\underline{\mathbf{B}}$ is contiguously accessed if i_q or j is incremented with unit-steps depending on the storage format of $\underline{\mathbf{B}}$. The access pattern can be improved by reordering tensor elements according to the storage format. However, copy operations reduce the overall throughput of the operation [15].

A better approach is to access tensor elements according to the tensor layout using the tensor layout tuple $\boldsymbol{\pi}$ as proposed in [4]. The modified algorithm 1 contiguously accesses memory for $\pi_1 \neq q$ and $p > 1$. Each recursion level adjusts only one multi-index element i_{π_r} with a stride w_{π_r} in line 5. With increasing recursion level and decreasing r , indices are incremented with smaller step sizes as $w_{\pi_r} \leq w_{\pi_{r+1}}$. The condition of the second **if** statement in line 4 is changed from $r \geq 1$ to $r > 1$. In this way, the mode- π_1 loop with index i_{π_1} and the minimum stride w_{π_1} are included in the base case which contains three loops performing a slice-matrix multiplication. The loop ordering are adjusted according to the tensor and matrix layout. The inner-most loop increments i_{π_1} and contiguously accesses tensor elements of $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$. The second loop increments i_q with which elements of $\underline{\mathbf{B}}$ are contiguously accessed if $\underline{\mathbf{B}}$ is stored in the row-major format. The third loop increments j and could be placed as the second loop if $\underline{\mathbf{B}}$ is stored in the column-major format.

```

1 tensor_times_matrix(A, B, C, n, i, m, q, q̂, r)
2   if  $r = \hat{q}$  then
3     | tensor_times_matrix(A, B, C, n, i, m, q, q̂,  $r - 1$ )
4   else if  $r > 1$  then
5     | for  $i_{\pi_r} \leftarrow 1$  to  $n_{\pi_r}$  do
6       | | tensor_times_matrix(A, B, C, n, i, m, q, q̂,  $r - 1$ )
7   else
8     | for  $j \leftarrow 1$  to  $m$  do
9       | | for  $i_q \leftarrow 1$  to  $n_q$  do
10        | | | for  $i_{\pi_1} \leftarrow 1$  to  $n_{\pi_1}$  do
11          | | | | C( $i_1, \dots, i_{q-1}, j, i_{q+1}, \dots, i_p$ ) += A( $i_1, \dots, i_q, \dots, i_p$ ) · B( $j, i_q$ )

```

Algorithm 1: Modified baseline algorithm with contiguous memory access for the tensor-matrix multiplication. The tensor order must be greater than one and for the contraction mode $1 \leq q \leq p$ and $\pi_1 \neq q$ must hold. The algorithm needs to be initially called with $r = p$ where \mathbf{n} is the shape tuple of $\underline{\mathbf{A}}$ and m is the q -th dimension of $\underline{\mathbf{C}}$.

While spatial data locality is improved by adjusting the loop ordering, the temporal data locality of tensors $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ differ. Note that slice $\underline{\mathbf{A}}'_{\pi_1, q}$ is accessed m times, fiber $\underline{\mathbf{C}}_{\pi_1}$ is accessed $\mathbf{n}(q)$ times and element $\mathbf{B}(j, i_q)$ is accessed $\mathbf{n}(\pi_1)$ times. While the specified fiber of $\underline{\mathbf{C}}$ can fit into first or second level cache, slice elements of $\underline{\mathbf{A}}$ are unlikely to fit in the local caches if the slice size $n_{\pi_1} \times n_q$ is large leading to higher cache misses and suboptimal performance. Optimized tiling for better temporal data locality has been discussed in [6] which suggests to use existing high-performance BLAS implementations for the base case.

4.3 BLAS-based Algorithms with Tensor Slices

Algorithm 1 is the starting point for the BLAS-based algorithm which computes the tensor-matrix product with a `gemm` routine. Besides the illustrated algorithm, we have identified seven other cases where a single `gemm` call suffices to compute the tensor-matrix product even if the tensor order $p > 2$. In summary, there are eight cases with a single `gemm` call using different arguments which are listed in table 1. The list of `gemm` calls supports all linear tensor layout and has no limitation on tensor order and contraction mode. The arguments of `gemm` are chosen depending on the tensor order p , tensor layout π and contraction mode q except for the `CBLAS_ORDER` which is `CblasRowMajor`.

Case 1 ($p = 1$): The tensor-vector product $\underline{\mathbf{A}} \times_1 \mathbf{B}$ can be computed with a `gemv` operation $\mathbf{a}^T \cdot \mathbf{B}$ where $\underline{\mathbf{A}}$ is an order-1 tensor, i.e. a vector \mathbf{a} of length n_1 .

Case 2-5 ($p = 2$): If $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ are order-2 tensors, i.e. a matrix \mathbf{A} with dimensions n_1 and n_2 , then a single `gemm` suffices to compute the tensor-matrix product. If \mathbf{A} and \mathbf{C} have the column-major format with $\pi = (1, 2)$, `gemm` either executes $\mathbf{C} = \mathbf{A} \cdot \mathbf{B}^T$ for $q = 1$ or $\mathbf{C} = \mathbf{B} \cdot \mathbf{A}$ for $q = 2$. Note that `gemm` interprets

Case	Order p	Layout π	Mode q	Routine	T	M	N	K	A	LDA	B	LDB	LDC
1	1	-	1	gemv	-	m	n_1	-	\mathbf{B}	n_1	$\underline{\mathbf{A}}$	-	-
2	2	(1, 2)	1	gemm	\mathbf{B}	n_2	m	n_1	$\underline{\mathbf{A}}$	n_1	\mathbf{B}	n_1	m
3	2	(1, 2)	2	gemm	-	m	n_1	n_2	\mathbf{B}	n_2	$\underline{\mathbf{A}}$	n_1	n_1
4	2	(2, 1)	1	gemm	-	m	n_2	n_1	\mathbf{B}	n_1	$\underline{\mathbf{A}}$	n_2	n_2
5	2	(2, 1)	2	gemm	\mathbf{B}	n_1	m	n_2	$\underline{\mathbf{A}}$	n_2	\mathbf{B}	n_2	m
6	> 2	any	π_1	gemm	\mathbf{B}	\bar{n}_q	m	n_q	$\underline{\mathbf{A}}$	n_q	\mathbf{B}	n_q	m
7	> 2	any	π_p	gemm	-	m	\bar{n}_q	n_q	\mathbf{B}	n_q	$\underline{\mathbf{A}}$	\bar{n}_q	\bar{n}_q
8	> 2	any	π_2, \dots, π_{p-1}	gemm*	-	m	n_{π_1}	n_q	\mathbf{B}	n_q	$\underline{\mathbf{A}}$	w_q	w_q

Table 1. Eight `gemv` and `gemm` cases for the mode- q tensor-matrix multiplication. Arguments T, M, N, etc. of the BLAS are chosen with respect to the tensor order p , layout π and contraction mode q where T specifies if \mathbf{B} is transposed. `gemm*` denotes multiple `gemm` calls with different tensor slices. Argument \bar{n}_q for case 6 and 7 is given by $\bar{n}_q = 1/n_q \prod_r n_r$. Matrix \mathbf{B} has the row-major format.

231 \mathbf{C} and \mathbf{A} as matrices using the reshaping operation ρ with $\rho = (2, 1)$ in row-
 232 major format even though both are stored column-wise. If \mathbf{A} and \mathbf{C} have the
 233 row-major format with $\pi = (2, 1)$, `gemm` either executes $\mathbf{C} = \mathbf{B} \cdot \mathbf{A}$ for $q = 1$ or
 234 $\mathbf{C} = \mathbf{A} \cdot \mathbf{B}^T$ for $q = 2$. Note that the transposition of \mathbf{B} is necessary for the cases
 235 2,5 and independent of the chosen storage format.

236 *Case 6-7 ($p > 2$):* If the order of $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ is greater than 2 and if the
 237 contraction mode q is equal to π_1 (case 6), a single `gemm` with the depicted
 238 parameters executes $\mathbf{C} = \mathbf{A} \cdot \mathbf{B}^T$ and computes a tensor-matrix product $\underline{\mathbf{C}} =$
 239 $\underline{\mathbf{A}} \times_{\pi_1} \mathbf{B}$ for any storage layout of $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$. Tensors $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ are flattened with
 240 $\varphi_{2,p}$ to row-major matrices \mathbf{A} and \mathbf{C} . Matrix \mathbf{A} has $\bar{n}_{\pi_1} = \bar{n}/n_{\pi_1}$ rows and n_{π_1}
 241 columns while matrix \mathbf{C} has the same number of rows and m columns. If $\pi_p = q$
 242 (case 7), Tensors $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ are flattened with $\varphi_{1,p-1}$ to column-major matrices
 243 \mathbf{A} and \mathbf{C} . Matrix \mathbf{A} has n_{π_p} rows and $\bar{n}_{\pi_p} = \bar{n}/n_{\pi_p}$ columns while matrix \mathbf{C}
 244 has m rows and the same number of columns. A single `gemm` executes $\mathbf{C} = \mathbf{B} \cdot \mathbf{A}$
 245 and computes the tensor-matrix product $\underline{\mathbf{C}} = \underline{\mathbf{A}} \times_{\pi_p} \mathbf{B}$ for any storage layout
 246 of $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$. Note that in all cases no copy operation is performed in order to
 247 compute the desired contraction, see subsection 3.

248 *Case 8 ($p > 2$):* If the tensor order is greater than 2 with $\pi_1 \neq q$ and $\pi_p \neq q$,
 249 the modified baseline algorithm 1 is used to successively call $\bar{n}/(n_q \cdot n_{\pi_1})$ times
 250 `gemm` with different tensor slices of $\underline{\mathbf{C}}$ and $\underline{\mathbf{A}}$ in the base case. Each `gemm` computes
 251 one slice $\underline{\mathbf{C}}'_{\pi_1,q}$ of the tensor-matrix product $\underline{\mathbf{C}}$ using the corresponding tensor
 252 slices $\underline{\mathbf{A}}'_{\pi_1,q}$ and the matrix \mathbf{B} . The matrix-matrix product $\mathbf{C} = \mathbf{B} \cdot \mathbf{A}$ is performed
 253 by interpreting both tensor slices as row-major matrices \mathbf{A} and \mathbf{C} which have
 254 the dimensions (n_q, n_{π_1}) and (m, n_{π_1}) , respectively.

255 4.4 BLAS-Based Algorithms with Subtensors

256 Case 8 can be optimized by selecting larger subtensors instead of tensor slices
 257 which might lead to a better processor utilization. Larger subtensors can be
 258 used by selecting additional mergeable modes that still allow the subtensor to be
 259 flattened into a matrix without reordering tensor elements using the description
 260 provided in section 3, see also lemma 4.1 in [10]. The maximum number of
 261 mergeable modes is $\hat{q} - 1$ with $\hat{q} = \pi^{-1}(q)$ and the corresponding modes are
 262 $\pi_1, \pi_2, \dots, \pi_{\hat{q}-1}$. Applying flattening $\varphi_{1,q-1}$ and reshaping ρ with $\boldsymbol{\rho} = (2, 1)$ on
 263 a subtensor of $\underline{\mathbf{A}}$ with dimensions $n_{\pi_1}, \dots, n_{\pi_{\hat{q}-1}}, n_q$ yields a row-major matrix
 264 \mathbf{A} with shape $(n_q, \prod_{r=1}^{\hat{q}-1} n_{\pi_r})$. This is done analogously for $\underline{\mathbf{C}}$ resulting in a row-
 265 major matrix with shape $(m, \prod_{r=1}^{\hat{q}-1} n_{\pi_r})$. This description supports all linear
 266 tensor layouts and generalizes lemma 4.2 in [10].

267 Algorithm 1 needs a minor modification for calling `gemm` flattened subtensors
 268 instead of tensor slices. The modified algorithm must therefor omit the first
 269 \hat{q} modes $\pi_{1,\hat{q}} = (\pi_1, \dots, \pi_{\hat{q}})$ including $\pi_{\hat{q}} = q$. This is done by only iterating
 270 over modes larger than \hat{q} in the non-base case. The conditions in line 2 and 4 are
 271 changed to $1 < r \leq \hat{q}$ and $\hat{q} < r$, respectively. The single indices of the subtensors
 272 $\underline{\mathbf{A}}'_{\pi_{1,\hat{q}}}$ and $\underline{\mathbf{C}}'_{\pi_{1,\hat{q}}}$ are given by the loop induction variables that belong to the
 273 π_r -th loop with $\hat{q} + 1 \leq r \leq p$.

274 4.5 Parallel BLAS-based Algorithms

275 The following paragraphs discuss three parallel approaches for the eighth case.
 276 Cases 1 to 7 already call a multi-threaded `gemm` and cannot be further optimized.

277 **Sequential Loops and Multithreaded `gemm`** One straight forward approach
 278 is to use algorithm 1 as it is and to sequentially call a multi-threaded `gemm` in
 279 the base case of the algorithm as described in subsection 4.4. This is beneficial if
 280 $q = \pi_{p-1}$, the inner dimensions n_{π_1}, \dots, n_q are large or the outer-most dimension
 281 n_{π_p} is smaller than the available processor cores. However, if the above conditions
 282 are not met, the processor cores might not be fully utilized where each multi-
 283 threaded `gemm` is executed with small subtensors. We will refer to this algorithm
 284 version as `<seq-loops,par-gemm>` that is executable with subtensors or tensor
 285 slices.
 286

287 **Parallel Loops and Single- or Multithreaded `gemm`** A more advanced version
 288 of the above algorithm executes a single-threaded `gemm` in parallel including
 289 all available (free) modes which depend on the slicing. If subtensors are used,
 290 all $\pi_{\hat{q}+1}, \dots, \pi_p$ modes are free. In case of tensor slices, only π_1 and $\pi_{\hat{q}}$ are free
 291 modes. The corresponding maximum degree of parallelism for both cases are
 292 $\prod_{r=\hat{q}+1}^p n_{\pi_r}$ and $\prod_{r=1}^p n_r / (n_{\pi_1} n_{\pi_{\hat{q}}})$, respectively.

293 Using tensor slices for the multiplication, $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ are flattened twice with
 294 $\varphi_{\pi_{\hat{q}+1}, \pi_p}$ and $\varphi_{\pi_2, \pi_{\hat{q}-1}}$. The resulting tensor is of order 4 with dimensions $n_{\pi_1},$
 295 $\hat{n}_{\pi_2}, n_q, \hat{n}_{\pi_4}$ where $\hat{n}_{\pi_2} = \prod_{r=2}^{\hat{q}-1} n_{\pi_r}$ and $\hat{n}_{\pi_4} = \prod_{r=\hat{q}+1}^p n_{\pi_r}$. In this way the tree-
 296 recursion has been transformed in two loops. The outer loop iterates over \hat{n}_{π_4}

while the inner loop iterates over \hat{n}_{π_2} calling `gemm` with slices $\underline{\mathbf{A}}'_{\pi_1,q}$ and $\underline{\mathbf{C}}'_{\pi_1,q}$. Both loops are parallelized using `omp parallel for` together with the `collapse(2)` and the `num_threads` clause which specifies the thread number.

In case of the general subtensor-matrix approach, both tensors are flattened twice with $\varphi_{\pi_{\hat{q}+1},\pi_p}$ and $\varphi_{\pi_1,\pi_{\hat{q}-1}}$. The resulting tensor is of order 3 with dimensions $\hat{n}_{\pi_1}, n_q, \hat{n}_{\pi_4}$ where $\hat{n}_{\pi_1} = \prod_{r=1}^{\hat{q}-1} n_{\pi_r}$ and $\hat{n}_{\pi_4} = \prod_{r=\hat{q}+1}^p n_{\pi_r}$. The corresponding algorithm consists of one loops which iterates over \hat{n}_{π_4} calling single-threaded `gemm` with multiple subtensors $\underline{\mathbf{A}}'_{\pi',q}$ and $\underline{\mathbf{C}}'_{\pi',q}$ with $\pi' = (\pi_1, \dots, \pi_{\hat{q}-1})$.

Both algorithm variants will be referred to as `<par-loops,seq-gemm>` which can be used with subtensors or tensor slices. Note that `<seq-loops,par-gemm>` and `<par-loops,seq-gemm>` are opposing versions where either `gemm` or the free loops are performed in parallel. The all-parallel version `<par-loops,par-gemm>` executes available loops in parallel where each loop thread executes a multi-threaded `gemm` with either subtensors or tensor slices.

Multithreaded `gemm_batch` The next version of the base algorithm is a modified version of the general subtensor-matrix approach that calls a single batched `gemm` for the eighth case. The subtensor dimensions and remaining `gemm` arguments remain the same. The library implementation is responsible how subtensor-matrix multiplications are executed and if subtensors are further divided into smaller subtensors or tensor slices. This version will be referred to as the `<gemm_batch>` variant.

5 Experimental Setup

Computing System The experiments have been carried out on an Intel Xeon Gold 6248R processor with a Cascade micro-architecture. The processor consists of 24 cores operating at a base frequency of 3 GHz for non-AVX512 instructions. With 24 cores and a peak AVX-512 boost frequency of 2.5 GHz, the processor achieves a theoretical data throughput of ca. 1.92 double precision TFlops/s. We measured a peak performance of 1.78 double precision Tflops/s using the `likwid` performance tool.

The source code has been compiled with `GCC v10.2` using the highest optimization level `-O3` and `-march=native`, `-pthread` and `-fopenmp`. Loops within for the eighth case have been parallelized using `GCC's OpenMP v4.5` implementation. We have used the `GEMV` and `GEMM` implementation of the 2024.0 Intel MKL and its own threading library `mk1_intel_thread` together with the threading runtime library `libiomp5`.

If not otherwise mentioned, both tensors $\underline{\mathbf{A}}$ and $\underline{\mathbf{C}}$ are stored according to the first-order linear tensor layout with $\pi = (1, \dots, p)$ whereas matrix \mathbf{B} has the row-major storage format. The benchmark results of each function are the average of 10 runs.

Tensor Shapes We have used asymmetrically-shaped and symmetrically-shaped tensors in order to cover many possible use cases. The dimension tuples of

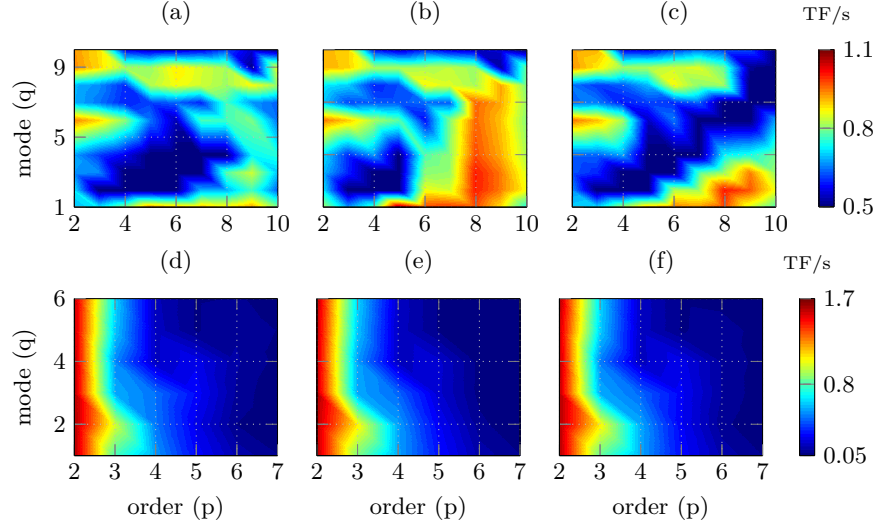


Fig. 1. Performance maps in double-precision Tflops/s of the proposed tensor-times-matrix algorithms with varying tensor orders p and contraction modes q . Tensors are asymmetrically-shaped on the upper plots and symmetrically-shaped on the lower plots. In (a) and (d) function `<gemm_batch>` is executed, in (b) and (e) `<par-loops,seq-gemm>` with tensor slices, in (c) and (f) `<par-loops,seq-gemm>` with subtensors.

both shape types are organized within two three-dimensional arrays with which tensors are initialized. The dimension array for the first shape type contains $720 = 9 \times 8 \times 10$ dimension tuples where the row number is the tensor order ranging from 2 to 10. For each tensor order 8 tensor instances with increasing tensor size is generated. The second set consists of $336 = 6 \times 8 \times 7$ dimensions tuples where the tensor order ranges from 2 to 7 and has 8 dimension tuples for each order. Each tensor dimension within the second set is 2^{12} , 2^8 , 2^6 , 2^5 , 2^4 and 2^3 . A detailed explanation of the tensor shape setup is given in [3, 4].

6 Results and Discussion

Slicing Methods The following paragraphs analyze the two proposed slicing methods by benchmarking the functions `<par-loops,seq-gemm>` and `<gemm_batch>` using asymmetrically (top) and symmetrically (bottom) shaped tensors. Fig. 1 contains six contour plots (performance maps) in which `<par-loops,seq-gemm>` either uses subtensors or tensor slices and `<gemm_batch>` loops over subtensors only. Each point within the performance map represents a mean value that has been averaged over tensor sizes for a tensor order¹.

¹ Note that Fig. 2 suggests that the contraction mode q can be greater than p which is not possible. We have kept the illustration for convenience, yet our profiling program sets $q = p$ in such cases.

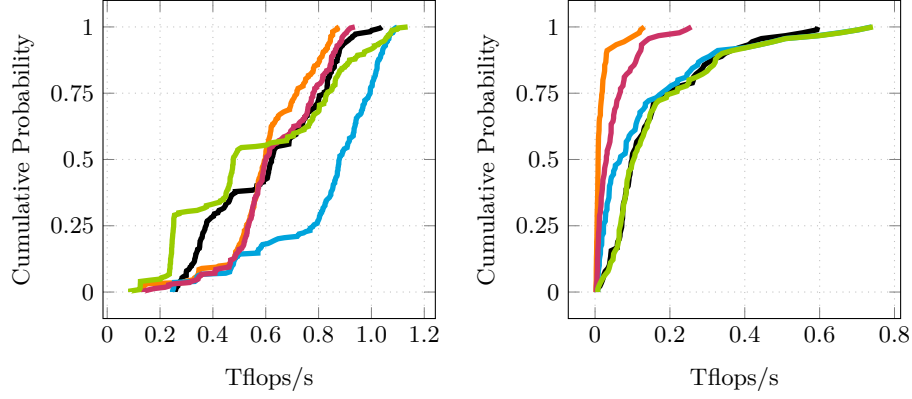


Fig. 2. Cumulative performance distributions of the proposed algorithms for case eight. Each distribution line belongs to one algorithm: `<gemm_batch>` (black), `<seq-loops,par-gemm>` (orange) and `<par-loops,seq-gemm>` (blue) using tensor slices, `<seq-loops,par-gemm>` (pink) and `<par-loops,seq-gemm>` (green) using subtenors. Tensors are asymmetrically (left plot) and symmetrically shaped (right plot).

For asymmetrically shaped tensors, function `<par-loops,seq-gemm>` with tensor slices performs on average 18% better than with subtenors. Our function `<par-loops,seq-gemm>` with tensor slices is on average 11% faster than Intel's `gemm_batch` routine and reaches almost 1.1 Tflops/s for non-edge cases with $q > 2$ and $p > 6$. This suggests that the Intel's implementation might not divide subtenors into smaller blocks.

With symmetrically shaped tensors, `<par-loops,seq-gemm>` with tensor slices performs almost identical as `<gemm_batch>` with 221.52 Gflops/s and 236.21 Gflops/s, respectively. Moreover, the slicing method almost has no affect on the runtime behavior of `<par-loops,seq-gemm>`. In contrast to the performance maps with asymmetrically shaped tensors, all functions almost reach the attainable peak performance of 1.7 Tflops/s when $p = 2$. This can be by the fact that both dimensions are equal or larger than 4096 enabling `gemm` to operate under optimal conditions.

Parallelization Methods The presented contour plots in Fig. 1 contain performance data of all cases except 4 and 5, see Table 1. The effects of the presented slicing and parallelization methods can be better understood if performance data of only the eighth case is examined. Fig. 2 contains cumulative performance distributions of all the proposed algorithms which are generated `gemm` or `gemm_batch` calls of case 8. As the distribution is empirically generated, the probability y of a point (x, y) on a distribution function corresponds to the number of test cases of a particular algorithm that achieves x Tflops/s or less. For instance, function `<seq-loops,par-gemm>` with subtenors computes the tensor-matrix product with equal to or less than 0.6 Tflops/s for 50% percent of the test cases using asymmetrically shaped tensor. Consequently, distribution functions with an ex-

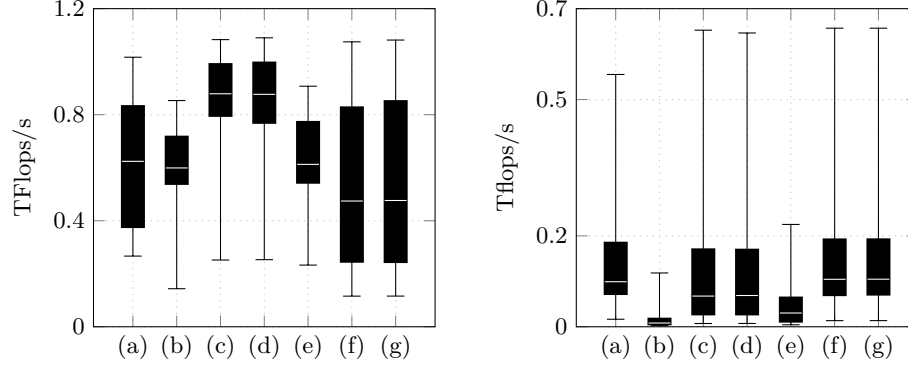


Fig. 3. Box plots visualizing performance statics in double-precision Tflops/s of a tensor-times-matrix algorithm for linear k -order tensor formats. The algorithm loops over single-threaded `gemm` with tensor slices with asymmetrically-shaped tensors on the left plot and with subtenors with symmetrically-shaped tensors on the right plot. Box plot number k denotes the utilized k -order storage.

ponential growth is favorable while logarithmic behavior is less desirable. The test set cardinality for case 8 is 255 for asymmetrically shaped tensors and 91 for symmetrically ones.

In case of asymmetrically shaped tensors, `<par-loops,seq-gemm>` with tensor slices performs best and outperforms `<gemm_batch>` and `<par-loops,seq-gemm>` with subtenors as previously stated. One unexpected finding is that function `<seq-loops,par-gemm>` with any slicing performs better than `<gemm_batch>` for certain tensor order and contraction mode combinations. All functions executed with symmetrically shaped tensors within case 8 only reach at most 743 Tflops/s instead of the attainable peak performance of 1.7 Tflops/s. This is expected as cases 2 and 3 for $p = 2$ with $q = 1$ and $q = 2$ are not considered. Functions `<par-loops,seq-gemm>` and `<gemm_batch>` have almost the same performance distribution outperforming `<seq-loops,par-gemm>` for almost every test case.

These findings

Comparison with other Approaches The following comparison includes three state-of-the-art libraries that implement three different approaches. The library `TCL (v0.1.1)` implements the (TTGT) approach with a high-perform tensor-transpose library `HPTT` which is discussed in [17]. `TBLIS (v1.0.0)` implements the `GETT` approach that is akin to `BLIS`'s algorithm design for matrix computations [12]. The tensor extension of `EIGEN (v3.3.90)` is used by the Tensorflow framework and performs the tensor-vector multiplication in-place and in parallel with contiguous memory access [1]. `TLIB` denotes our library that consists of sequential and parallel versions of the tensor-vector multiplication. Numerical results of `TLIB` have been verified with the ones of `TCL`, `TBLIS` and `EIGEN`.

Fig. ?? illustrates the average single-precision Gflops/s with asymmetrically- and symmetrically-shaped tensors in the first-order storage format. The run-

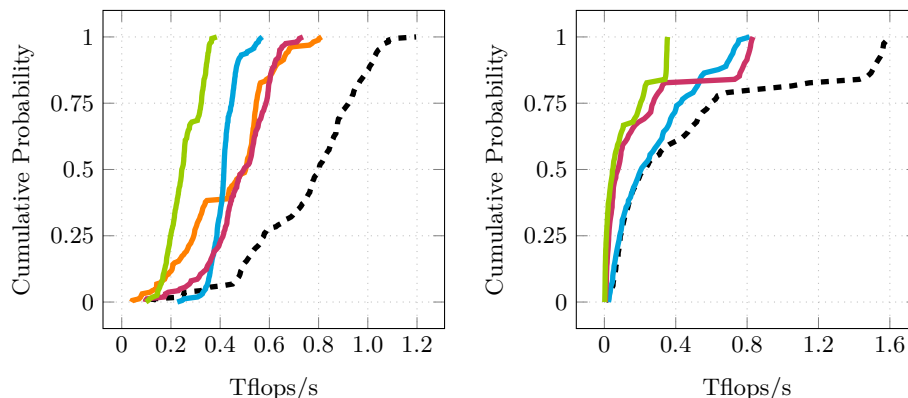


Fig. 4. Cumulative performance distributions of tensor-times-matrix algorithms in double-precision Tflops/s. Each distribution line belongs to a library: **tlib** (---), **tcl** (—), **tblis** (—), **libtorch** (—), **eigen** (—). Libraries have been tested with asymmetrically-shaped (left plot) and symmetrically-shaped tensors (right plot).

time behavior of **TBLIS** and **EIGEN** with asymmetrically-shaped tensors is almost constant for varying tensor sizes with a standard deviation ranging between 2% and 13%. **TCL** shows a different behavior with 2 and 4 Gflops/s for any order $p \geq 2$ peaking at $p = 10$ and $q = 2$. The performance values however deviate from the mean value up to 60%. Computing the arithmetic mean over the set of contraction modes yields a standard deviation of less than 10% where the performance increases with increasing order peaking at $p = 10$. **TBLIS** performs best for larger contraction dimensions achieving up to 7 Gflops/s and slower runtimes with decreasing contraction dimensions. In case of symmetrically-shaped tensors, **TBLIS** and **TCL** achieve up to 12 and 25 Gflops/s in single precision with a standard deviation between 6% and 20%, respectively. **TCL** and **TBLIS** behave similarly and perform better with increasing contraction dimensions. **EIGEN** executes faster with decreasing order and increasing contraction mode with at most 8 Gflops/s at $p = 2$ and $q \geq 2$.

Fig. ?? illustrates relative performance maps of the same tensor-vector multiplication implementations. Comparing **TCL** performance, **TLIB-SB-PN** achieves an average speedup of 6x and more than 8x for 42% of the test cases with asymmetrically shaped tensors and executes on average 5x faster with symmetrically shaped tensors. In comparison with **TBLIS**, **TLIB-SB-PN** computes the tensor-vector product on average 4x and 3.5x faster for asymmetrically and symmetrically shaped tensors, respectively.

7 Conclusion and Future Work

Based on the **LOG** approach, we have presented in-place and parallel tensor-vector multiplication algorithms of **TLIB**. Using highly-optimized **DOT** and **GEMV** routines

of `OpenBLAS`, our proposed algorithm is designed for dense tensors with arbitrary order, dimensions and any non-hierarchical storage format. `TLIB`'s algorithms either directly call `DOT`, `GEMV` or recursively perform parallel slice-vector multiplications using `GEMV` with tensor slices and fibers.

Our findings show that loop-fusion improves the performance of `TLIB`'s parallel version on average by a factor of 5x achieving up to 34.8/15.5 Gflops/s in single/double precision for asymmetrically shaped tensors. With symmetrically shaped tensors resulting in small contraction dimensions, the results suggest that higher-order slices with larger dimensions should be used. We have demonstrated that the proposed algorithms compute the tensor-vector product on average 6.1x and up to 12.6x faster than the `TTGT`-based implementation provided by `TCL`. In comparison with `TBLIS`, `TLIB` achieves speedups on average of 4.0x and at most 10.4x. In summary, we have shown that a `LOG`-based tensor-vector multiplication implementation can outperform current implementations that use a `TTGT` and `GETT` approaches.

In the future, we intend to design and implement the tensor-matrix multiplication with the same requirements also supporting tensor transposition and subtensors. Moreover, we would like to provide an in-depth analysis of `LOG`-based implementations of tensor contractions with higher arithmetic intensity.

Project and Source Code Availability `TLIB` has evolved from the Google Summer of Code 2018 project for extending `Boost`'s `uBLAS` library with tensors. Project description and source code can be found at <https://github.com/basso/ttv>. The sequential tensor-vector multiplication of `TLIB` is part of `uBLAS` and in the official release of `Boost v1.70.0`.

Acknowledgements The author would like to thank Volker Schatz and Banu Sözüar for proofreading. He also thanks Michael Arens for his support.

References

1. Abadi, M., Barham, P., Chen, J., Chen, Z., Davis, A., Dean, J., Devin, M., ..., Zheng, X.: Tensorflow: A system for large-scale machine learning. In: Proceedings of the 12th USENIX Conference on Operating Systems Design and Implementation. pp. 265–283. OSDI'16, USENIX Association, Berkeley, CA, USA (2016)
2. Bader, B.W., Kolda, T.G.: Algorithm 862: Matlab tensor classes for fast algorithm prototyping. *ACM Trans. Math. Softw.* **32**, 635–653 (December 2006)
3. Basso, C.: Design of a high-performance tensor-vector multiplication with blas. In: International Conference on Computational Science. pp. 32–45. Springer (2019)
4. Basso, C., Schatz, V.: Fast higher-order functions for tensor calculus with tensors and subtensors. In: International Conference on Computational Science. pp. 639–652. Springer (2018)
5. Golub, G.H., Van Loan, C.F.: Matrix Computations. JHU Press, 4 edn. (2013)
6. Goto, K., Geijn, R.A.v.d.: Anatomy of high-performance matrix multiplication. *ACM Transactions on Mathematical Software (TOMS)* **34**(3) (2008)

- 470 7. Karahan, E., Rojas-López, P.A., Bringas-Vega, M.L., Valdés-Hernández, P.A.,
471 Valdes-Sosa, P.A.: Tensor analysis and fusion of multimodal brain images. Pro-
472 ceedings of the IEEE **103**(9), 1531–1559 (2015)
- 473 8. Kolda, T.G., Bader, B.W.: Tensor decompositions and applications. SIAM review
474 **51**(3), 455–500 (2009)
- 475 9. Lee, N., Cichocki, A.: Fundamental tensor operations for large-scale data analysis
476 using tensor network formats. Multidimensional Systems and Signal Processing
477 **29**(3), 921–960 (2018)
- 478 10. Li, J., Battaglino, C., Perros, I., Sun, J., Vuduc, R.: An input-adaptive and in-place
479 approach to dense tensor-times-matrix multiply. In: High Performance Computing,
480 Networking, Storage and Analysis, 2015 SC-International Conference for. pp. 1–12.
481 IEEE (2015)
- 482 11. Lim, L.H.: Tensors and hypermatrices. In: Hogben, L. (ed.) Handbook of Linear
483 Algebra. Chapman and Hall, 2 edn. (2017)
- 484 12. Matthews, D.A.: High-performance tensor contraction without transposition.
485 SIAM Journal on Scientific Computing **40**(1), C1–C24 (2018)
- 486 13. Napoli, E.D., Fabregat-Traver, D., Quintana-Ortí, G., Bientinesi, P.: Towards an
487 efficient use of the blas library for multilinear tensor contractions. Applied Math-
488 ematics and Computation **235**, 454 – 468 (2014)
- 489 14. Papalexakis, E.E., Faloutsos, C., Sidiropoulos, N.D.: Tensors for data mining and
490 data fusion: Models, applications, and scalable algorithms. ACM Transactions on
491 Intelligent Systems and Technology (TIST) **8**(2), 16 (2017)
- 492 15. Shi, Y., Niranjana, U.N., Anandkumar, A., Cecka, C.: Tensor contractions with
493 extended blas kernels on cpu and gpu. In: 2016 IEEE 23rd International Conference
494 on High Performance Computing (HiPC). pp. 193–202 (Dec 2016)
- 495 16. Solomonik, E., Matthews, D., Hammond, J., Demmel, J.: Cyclops tensor frame-
496 work: Reducing communication and eliminating load imbalance in massively par-
497 allel contractions. In: Parallel & Distributed Processing (IPDPS), 2013 IEEE 27th
498 International Symposium on. pp. 813–824. IEEE (2013)
- 499 17. Springer, P., Bientinesi, P.: Design of a high-performance gemm-like tensor–tensor
500 multiplication. ACM Transactions on Mathematical Software (TOMS) **44**(3), 28
501 (2018)